

Understanding the Problem and RL Proposition

My recent capstone idea focuses on reducing the 56% tomato loss in Rwanda's supply chain from Farm to Collection Point, then to Wholesale, and finally to Retail. The goal is to minimize loss and boost profit. I had the intuition to use a Reinforcement Learning (RL) algorithm, which could learn the most effective ways to manage tomatoes based on the current state, including factors like location, weather, ripeness, and more. These are just a few of the variables I've been exploring. I believe RL could enhance the supply chain by identifying patterns of loss from actions being taken and preventing them. RL might suggest context-aware actions, but I suspect it needs help in understanding the root causes of losses, which is where Hidden Markov Models (HMMs) could assist.

Applying HMM to Support RL

I envisioned using HMMs to analyze loss percentages between transitions (e.g., [25%, 10%, 5%], where 25% is lost from Farm to Collection Point) to uncover hidden patterns to help us understand the causes of these losses which could be various factors but in the form of probabilities. With the Baum-Welch algorithm, HMMs could link these losses to hidden states such as "Intact" (low loss, e.g., 5%), "At Risk" (moderate, e.g., 10%), and "Compromised" (high, e.g., 25%). Research shows that factors like weather, ripeness, storage, and handling contribute to loss variability, factors that a human might overlook. I believe the HMM-RL combination enhances human decision-making: HMMs detect underlying patterns and predict upcoming losses (e.g., 5% after a 10% loss), while RL proposes tailored actions.

For instance, if HMM infers a "Compromised" state with a 25% loss due to early damage, RL could recommend actions appropriate from the Farm stage which would help maximize rewards. I think this collaborative approach could significantly reduce the 56% loss rate.

- **Farm to Collection Point:** Actions are guided by RL based on HMM insights.
- **Collection Point to Wholesale:** Actions are adjusted by RL using HMM insights.
- **Wholesale to Retail:** Actions are informed by RL with support from HMM insights.

Answering the Question: HMM Problem and Details

- **Observations:** The observed data consists of loss percentages (e.g., 25%, 10%) representing tomato losses at each supply chain transition.
- **Type of HMM Problem:** Since the hidden states ("Intact," "At Risk," "Compromised"), and the model parameters are not known beforehand and must be inferred from data, this is a **Baum-Welch problem**.
- **Algorithm to Train HMM:** The **Baum-Welch algorithm** is used to train the HMM based on the observed loss data.
- **Known Values:** The known values are the observed losses at each stage in the supply chain.
- **Unknown Values:** The unknowns include the hidden states and the predicted future losses, and the probabilities, which are learned over time.

- **Parameters Updated:** Parameters such as transition probabilities (between hidden states), emission probabilities (likelihood of observed losses given a state), and initial state probabilities are updated during training.